

QUESTION 1

Question 1

Display the data types of each column using the function dtypes. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[14]: #Enter Your Code, Execute and take the Screenshot
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 22 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Unnamed: 0           21613 non-null  int64
1   id                   21613 non-null  int64
2   date                 21613 non-null  object
3   price                21613 non-null  float64
4   bedrooms             21600 non-null  float64
5   bathrooms            21603 non-null  float64
6   sqft_living          21613 non-null  int64
7   sqft_lot             21613 non-null  int64
8   floors               21613 non-null  float64
9   waterfront           21613 non-null  int64
10  view                 21613 non-null  int64
11  condition            21613 non-null  int64
12  grade                21613 non-null  int64
13  sqft_above           21613 non-null  int64
14  sqft_baseament       21613 non-null  int64
15  yr_built              21613 non-null  int64
16  yr_renovated         21613 non-null  int64
17  zipcode              21613 non-null  int64
18  lat                  21613 non-null  float64
19  long                 21613 non-null  float64
20  sqft_living15        21613 non-null  int64
21  sqft_lot15           21613 non-null  int64
dtypes: float64(6), int64(15), object(1)
memory usage: 3.5+ MB
```

QUESTION 2

Module 2: Data Wrangling

Question 2

Drop the columns "id" and "Unnamed: 0" from axis 1 using the method `drop()`, then use the method `describe()` to obtain a statistical summary of the data. Make sure the `inplace` parameter is set to `True`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
#Enter Your Code, Execute and take the Screenshot
df = df.drop(columns=['id','Unnamed: 0'])
df
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_baseament	yr_built	yr_renovated	zipcode	lat	long	sqft_livi
0	20141013T000000	221900.0	3.0	1.00	1180	5650	1.0	0	0	3	7	1180	0	1955	0	98178	47.5112	-122.257	
1	20141209T000000	538000.0	3.0	2.25	2570	7242	2.0	0	0	3	7	2170	400	1951	1991	98125	47.7210	-122.319	
2	20150225T000000	180000.0	2.0	1.00	770	10000	1.0	0	0	3	6	770	0	1933	0	98028	47.7379	-122.233	
3	20141209T000000	604000.0	4.0	3.00	1960	5000	1.0	0	0	5	7	1050	910	1965	0	98136	47.5208	-122.393	
4	20150218T000000	510000.0	3.0	2.00	1680	8080	1.0	0	0	3	8	1680	0	1987	0	98074	47.6168	-122.045	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
21608	20140521T000000	360000.0	3.0	2.50	1530	1131	3.0	0	0	3	8	1530	0	2009	0	98103	47.6993	-122.346	
21609	20150223T000000	400000.0	4.0	2.50	2310	5813	2.0	0	0	3	8	2310	0	2014	0	98146	47.5107	-122.362	
21610	20140623T000000	402101.0	2.0	0.75	1020	1350	2.0	0	0	3	7	1020	0	2009	0	98144	47.5944	-122.299	
21611	20150116T000000	400000.0	3.0	2.50	1600	2388	2.0	0	0	3	8	1600	0	2004	0	98027	47.5345	-122.069	
21612	20141015T000000	325000.0	2.0	0.75	1020	1076	2.0	0	0	3	7	1020	0	2008	0	98144	47.5941	-122.229	

21613 rows × 20 columns

## QUESTION 3

### Module 3: Exploratory Data Analysis

#### Question 3

Use the method `value_counts()` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a data frame. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[26]: #Enter Your Code, Execute and take the Screenshot
df['floors'].value_counts().to_frame()
```

```
[26]:
```


	count
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

## QUESTION 4

#### Question 4

Use the function `boxplot` in the seaborn library to determine whether houses with a waterfront view or without a waterfront view have more price outliers. Take a screenshot of your code and boxplot. You will need to submit the screenshot for the final project.

```
[31]: sns.boxplot(x="waterfront",y = "price", data = df)
xlabel = "waterfront"
ylabel = "price"
```




## QUESTION 5

### Question 5

Use the function `regplot` in the seaborn library to determine if the feature `sqft_above` is negatively or positively correlated with price. Take a screenshot of your code and scatterplot. You will need to submit the screenshot for the final project.

```
[32]: #Enter Your Code, Execute and take the Screenshot
sns.regplot(x="sqft_above",y = "price", data = df)
plt.grid()
```



## QUESTION 6

### Module 4: Model Development

We can Fit a linear regression model using the longitude feature `'long'` and calculate the  $R^2$ .

```
[37]: X = df[['long']]
Y = df['price']
lm = LinearRegression()
lm.fit(X,Y)
lm.score(X, Y)
```

```
[37]: 0.00046769430149007363
```

### Question 6

Fit a linear regression model to predict the `'price'` using the feature `'sqft_living'` then calculate the  $R^2$ . Take a screenshot of your code and the value of the  $R^2$ . You will need to submit it for the final project.

```
[38]: #Enter Your Code, Execute and take the Screenshot
X = df[['sqft_living']]
Y = df['price']
lm = LinearRegression()
lm.fit(X,Y)
lm.score(X, Y)
```

```
[38]: 0.4928532179037931
```

## QUESTION 7

### Question 7

Fit a linear regression model to predict the `'price'` using the list of features:

```
[53]: Z =df[['floors', "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view" ,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]]
```

Then calculate the  $R^2$ . Take a screenshot of your code and the value of the  $R^2$ . You will need to submit it for the final project.

```
[54]: #Enter Your Code, Execute and take the Screenshot
lm.fit(Z,Y)
print(lm.score(Z, Y))
```

```
0.6576890354915759
```

## QUESTION 8

### This will help with Question 8

Create a list of tuples, the first element in the tuple contains the name of the estimator:

```
'scale'
```

```
'polynomial'
```

```
'model'
```

The second element in the tuple contains the model constructor

```
StandardScaler()
```

```
PolynomialFeatures(include_bias=False)
```

```
LinearRegression()
```

```
[55]: Input=[('scale',StandardScaler()),('polynomial', PolynomialFeatures(include_bias=False)),('model',LinearRegression())]
```

### Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list `features`, and calculate the  $R^2$ . Take a screenshot of your code and the value of the  $R^2$ . You will need to submit it for the final project.

```
[61]: #Enter Your Code, Execute and take the Screenshot
```

```
pipe=Pipeline(Input)
Z = Z.astype(float)
pipe.fit(Z,Y)
ypipe=pipe.predict(Z)
R2 = pipe.score(Z, Y)
print(f'R^2: {R2}')

R^2: 0.7512051345272872
```

## QUESTION 9

### Module 5: Model Evaluation and Refinement

Import the necessary modules:

```
[26]: from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
print("done")
```

done

We will split the data into training and testing sets:

```
[27]: features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view" ,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
X = df[features]
Y = df['price']

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.15, random_state=1)

print("number of test samples:", x_test.shape[0])
print("number of training samples:",x_train.shape[0])

number of test samples: 3242
number of training samples: 18371
```

### Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the  $R^2$  using the test data. Take a screenshot of your code and the value of the  $R^2$ . You will need to submit it for the final project.

```
[28]: from sklearn.linear_model import Ridge
```

```
[29]: #Enter Your Code, Execute and take the Screenshot
```

```
ridge = Ridge(alpha=0.1)
ridge.fit(x_train, y_train)
y_pred = ridge.predict(x_test)
R2_test = ridge.score(x_test, y_test)
print(f'R^2 on test data: {R2_test}')
```

R^2 on test data: 0.647875916393907

## QUESTION 10

### Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the  $R^2$  utilising the test data provided. Take a screenshot of your code and the  $R^2$ . You will need to submit it for the final project.

```
[30]: #Enter Your Code, Execute and take the Screenshot
poly = PolynomialFeatures(degree=2, include_bias=False)
x_train_poly = poly.fit_transform(x_train)
x_test_poly = poly.transform(x_test)
ridge_model = Ridge(alpha=0.1)
ridge_model.fit(x_train_poly, y_train)
y_predpoly = ridge_model.predict(x_test_poly)
R2_test = ridge_model.score(x_test_poly, y_test)
print(f'R^2 on test data: {R2_test}')
```

R^2 on test data: 0.700274425803224